

Python For Finance

M1 - Economie & Finance - Gradient Descent

Université Paris Dauphine

October 2025

Part I — Definition

For a function $f(x_1, \dots, x_n)$, the gradient is the vector of partial derivatives:

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n} \right).$$

The gradient indicates the direction of the steepest local increase.

Its norm, $\|\nabla f\|$, is the local rate of increase (its magnitude / intensity).


Part I — Definition (example)

Let $f(x, y) = x^2 + y^2$. Its gradient is

$$\nabla f(x, y) = (2x, 2y).$$

At the point $(x, y) = (1, 2)$:

$$\nabla f(1, 2) = (2 \cdot 1, 2 \cdot 2) = (2, 4), \quad \|\nabla f(1, 2)\| = \|(2, 4)\| = \sqrt{2^2 + 4^2} = \sqrt{20}.$$



Surface $z = x^2 + y^2$ (gradient at $(1, 2)$ has direction $(2, 4)$ and magnitude $\sqrt{20}$)

Part II — Gradient Descent

Gradient descent is an iterative optimization method for minimizing a function $f(\theta)$.

At each step, we move the parameters θ in the direction opposite to the gradient (the steepest-ascent direction) because that yields the steepest decrease of f .

Typical update:

$$\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} f(\theta_t)$$

where $\alpha > 0$ is the learning rate and $\nabla_{\theta} f$ the gradient.

Part II — Multivariate example

Suppose we have a dataset giving the living areas, the number of bedrooms and prices for houses:

Living area (ft ²)	#bedrooms	Price (1000\$)
2104	3	400
1600	3	330
2400	3	369
1416	2	232
3000	4	540
:	:	:

Here, the x 's are two-dimensional vectors in \mathbb{R}^2 . For instance, $x_1^{(i)}$ is the living area of the i -th house, and $x_2^{(i)}$ is its number of bedrooms.

Part II — Hypothesis and notation

Let's approximate y (house price) as a linear function of x (living area and nb of bedrooms):

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2.$$

Here, the θ_i are the **parameters** (also called *weights*).

To simplify notation, we drop the θ subscript in $h_{\theta}(x)$ and simply write $h(x)$ and we also introduce the convention $x_0 = 1$ (the *intercept* term), so that

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^\top x,$$

where on the right-hand side we view θ and x as vectors, and n is the number of input variables (not counting x_0).

Part II — Choosing the parameters θ

How do we pick, or learn, the parameters θ ?

One reasonable method is to make $h(x)$ close to y , at least for the training examples we have. To formalize this, define a function that measures, for each value of the θ 's, how close the $h(x^{(i)})$'s are to the corresponding $y^{(i)}$'s. We define the **cost function**:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2.$$

With m , the number of data point in the dataset. You may recognize this as the familiar *least-squares* cost function of the ordinary least squares (OLS) regression model. Indeed, If the matrix $X^\top X$ is invertible, the unique solution is

$$\theta = (X^\top X)^{-1} X^\top y.$$

Part II — Gradient Descent Update

We want to choose θ to minimize $J(\theta)$.

We use a search algorithm that starts from an initial guess and repeatedly changes θ to make $J(\theta)$ smaller, hoping to converge to a minima => **gradient descent**

The gradient descent performs the update simultaneously for all $j = 0, \dots, n$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta).$$

Here, α is the **learning rate**. This algorithm repeatedly steps in the direction of steepest decrease of J .

Part II — Gradient descent update

To implement the algorithm, we need $\frac{\partial}{\partial \theta_j} J(\theta)$.

For a single data point (x, y) , we have:

$$J(\theta) = \frac{1}{2} (h_{\theta}(x) - y)^2$$

with
$$h_{\theta}(x) = \sum_{i=0}^n \theta_i x_i,$$

Thus

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\ &= (h_{\theta}(x) - y) \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y) \\ &= (h_{\theta}(x) - y) x_j.\end{aligned}$$

Therefore, for a single example $(x^{(i)}, y^{(i)})$ the update rule becomes

$$\boxed{\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}}$$

Part II — LMS Update

The rule is called the **LMS** update rule (Least Mean Squares), also known as the **Widrow–Hoff** learning rule.

Per-example update (simultaneous for all j):

$$\theta_j \leftarrow \theta_j + \alpha(y^{(i)} - h_\theta(x^{(i)})) x_j^{(i)}$$


Vector form:

$$\theta \leftarrow \theta + \alpha(y^{(i)} - h_\theta(x^{(i)})) x^{(i)}.$$

Why it's intuitive:

- The update size is *proportional to the error* $(y^{(i)} - h_\theta(x^{(i)}))$.
- If the prediction nearly matches $y^{(i)}$, the change is small; if the error is large, the change is larger.

Part II — LMS Update



trajectory taken by gradient descent

Part I — Batch Gradient Descent (update rule)

Repeat until convergence:

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_\theta(x^{(i)})) x_j^{(i)} \quad \text{for every } j = 0, \dots, n$$

Batch gradient descent:

- Uses the *entire* data set each step
- The update direction is the negative gradient of the squared-error cost.
- Training large models is computationally intensive.

Part II — Reduce Compute: Use Stochastic Gradient Descent (SGD)

Motivation. Full-batch gradient descent processes all m examples each step \Rightarrow computationally intensive. To reduce compute per update, use **SGD**: update with one example at a time.

SGD update (squared error), per example $(x^{(i)}, y^{(i)})$:

$$\theta \leftarrow \theta + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x^{(i)}$$

(Component-wise: $\theta_j \leftarrow \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$.)

How to run it (presentation):

- Shuffle the training set.
- For each example $(x^{(i)}, y^{(i)})$: compute the error $y^{(i)} - h_{\theta}(x^{(i)})$ and apply the update.
- Repeat for multiple epochs; optionally decay α .

Part III — Why do we use Gradient descent

Why linear regression with the *least-squares* cost function is a natural choice under some assumptions.

Assumptions.

- Data follow a linear model with noise: $y^{(i)} = \theta^\top x^{(i)} + \varepsilon^{(i)}$.
- Errors are IID Gaussian: $\varepsilon^{(i)} \sim \mathcal{N}(0, \sigma^2)$.

Implication. We can write the distribution of $y(i)$ given $x(i)$ and parameterized by θ as follows

$$p(y^{(i)} | x^{(i)}; \theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2}\right)$$

Thus,

$$y^{(i)} | x^{(i)}; \theta \sim \mathcal{N}(\theta^\top x^{(i)}, \sigma^2)$$

Part III — Likelihood of the Data

Given X (the matrix collecting all $x^{(i)}$) and parameters θ , the distribution of the targets $\tilde{y} = (y^{(1)}, \dots, y^{(m)})$ is $p(\tilde{y} | X; \theta)$.

This is typically viewed as a function of \tilde{y} expressed given X for fixed θ . When we view it as a function of θ , it is called the **likelihood**:

$$L(\theta) = L(\theta; X, \tilde{y}) = p(\tilde{y} | X; \theta).$$

Assuming the errors $\varepsilon^{(i)}$ are IID Gaussian, we have

$$L(\theta) = \prod_{i=1}^m p(y^{(i)} | x^{(i)}; \theta) = \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2}\right).$$

Maximum likelihood principle: choose θ to make the observed data most probable, i.e.,

$$\theta^* \in \arg \max_{\theta} L(\theta).$$

Part III — Maximizing the Log-Likelihood

In practice we maximize the **log-likelihood** $\ell(\theta)$:

$$\ell(\theta) = \log L(\theta) = \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2}\right).$$

$$\ell(\theta) = \sum_{i=1}^m \left[\log \frac{1}{\sqrt{2\pi}\sigma} - \frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2} \right]$$

$$\ell(\theta) = m \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^m (y^{(i)} - \theta^\top x^{(i)})^2.$$

Part III — From MLE to Least Squares

Maximum likelihood principle: choose θ that maximizes $\ell(\theta)$.

Observation. The first term $m \log(1/(\sqrt{2\pi}\sigma))$ does not depend on θ . Therefore,

$$\arg \max_{\theta} \ell(\theta) = \arg \min_{\theta} \sum_{i=1}^m (y^{(i)} - \theta^\top x^{(i)})^2.$$

Equivalence (up to a constant factor):

$$\boxed{\arg \min_{\theta} \frac{1}{2} \sum_{i=1}^m (y^{(i)} - \theta^\top x^{(i)})^2 = \arg \max_{\theta} \ell(\theta)}$$

Hence, **least-squares regression = maximum likelihood** estimation under IID Gaussian noise.

Part I — Notes and Consequences

- σ^2 **not needed to find θ :** the MLE of θ does not depend on the (unknown) noise variance σ^2 .
- **Interpretation:** least squares minimizes the sum of squared residuals, which matches the negative log-likelihood under Gaussian noise.
- **Beyond Gaussians:** other noise models \Rightarrow other loss functions (e.g., Laplace $\Rightarrow \ell_1$ loss; Bernoulli \Rightarrow logistic loss).
- **Takeaway:** under simple, common assumptions, least squares is a *natural* and *principled* choice via maximum likelihood.

Part IV — Application

Now let's implement code to test the concepts explained above.

Follow the instructions in the Jupyter notebook titled *Gradient Descent*.

For simulating samples from a specified distribution, use the NumPy package.

To plot your results, use Matplotlib. If you need help, consult the official documentation.

Part IV — Progress Bars with tqdm

tqdm is a Python library that provides progress bars. Wrap any iterable with `tqdm(iterable)` to display live progress while your loop runs.

Example):

```
from tqdm import tqdm
import time

for i in tqdm(range(100), desc="Working"):
    time.sleep(0.01) # your work here
```